

## Routing

— So far discussed about packet forwarding

— Today : Routing , i.e.

How is that forwarding table created?

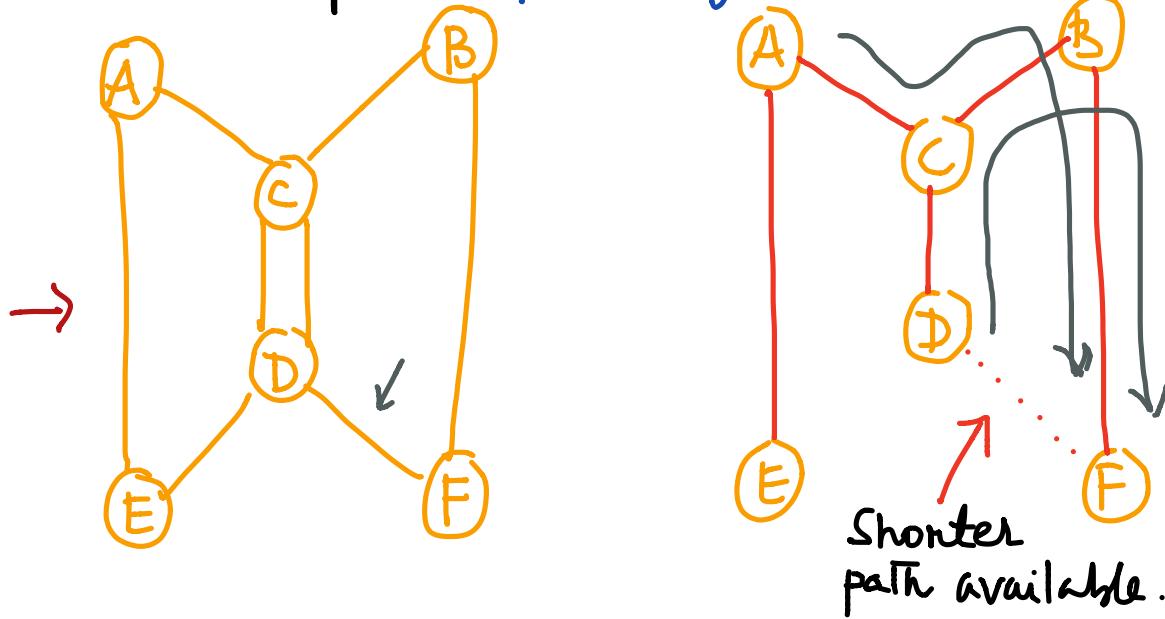


Forwarding is a local process.

Routing is a global process - needs agreement among all nodes that are involved in it.

---

Past Example: Spanning Tree

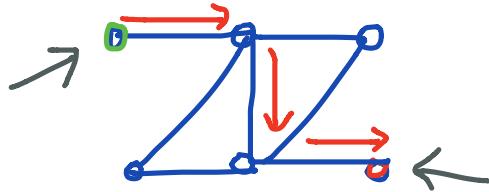


Obs: Spanning Tree algorithm provided  
a route, not necessarily an efficient  
route

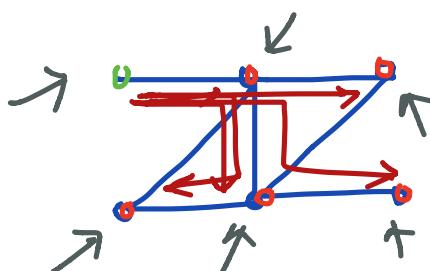
- Our goal in routing is to find "best" paths - e.g., least costly path.

### Delivery models :

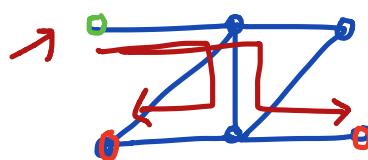
~~①~~ **Unicast**



~~②~~ **Broadcast**

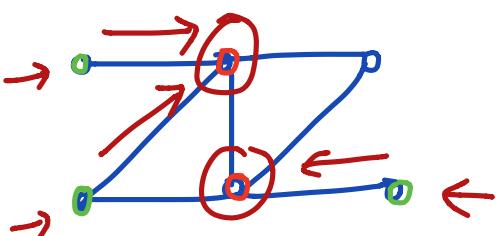


~~③~~ **Multi cast**



~~④~~ **Anycast**

e.g., DNS, CDN



## Objectives of Routing Algorithms

- ✓ ① Correctness - from source to destination
  - ✓ ② Efficient path finding - least cost
  - ✗ ③ Fast convergence - recovers quickly after topology changes
  - ✓ ④ Scalability - complexity of the algorithm not too high.
- 

Usually,

Routing is done in a **decentralized, distributed setting**

[these two are independent properties]

- ✓ no central controller
  - nodes can locally exchange messages with their neighbors
- ✓ nodes work in parallel
- ✓ they are susceptible to failures

## Topics to cover

### - Shortest Path Routing —

#### - Dijkstra's algorithm

✓ Distance Vector Routing

✓ Link-state routing

✓ Equal-cost multipath

✓ Inter-domain routing

- Border Gateway Protocol (BGP)

---

## Shortest Path Routing

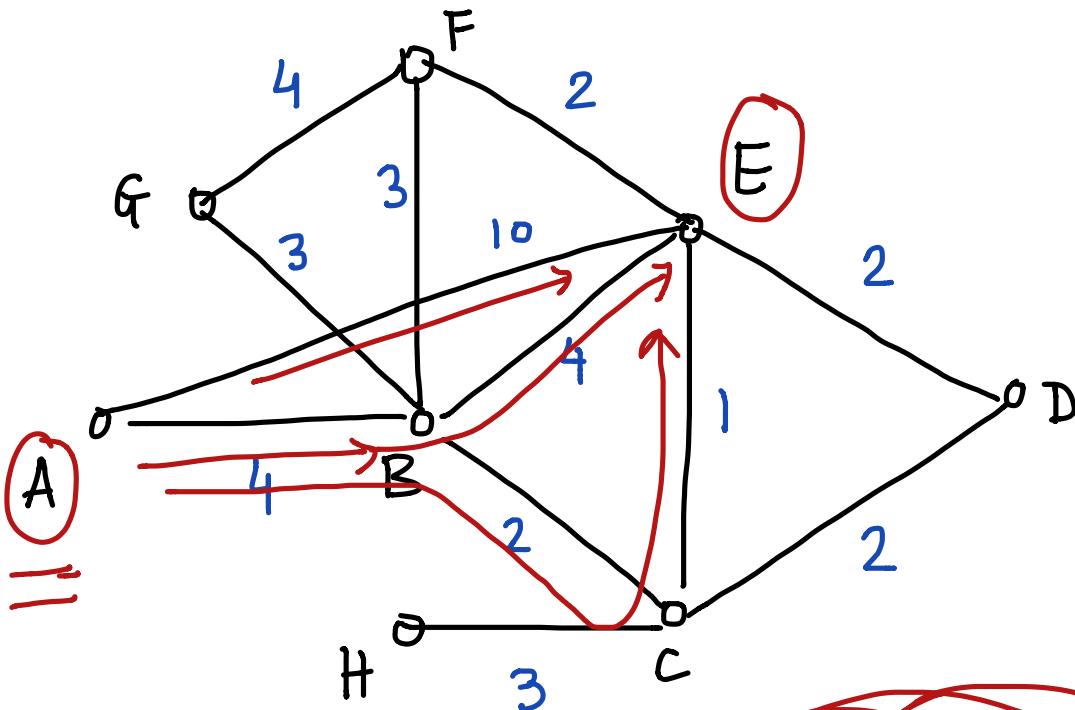
- one way of quantifying the "efficient" path

- "cost" reflects lot of factors, e.g., latency, expense of a link.

★ - but only static factors topology is fixed  
- we won't consider dynamic factors like the load of a link.

⊖ costs are additive over links ↴ ↓  
 $\text{cost}(A \rightarrow B \rightarrow C) = \text{cost}(A \rightarrow B) + \text{cost}(B \rightarrow C)$

Ex.



- Find shortest path  $A \rightarrow E$

- all links are bidirectional.

$A \rightarrow B \rightarrow C \rightarrow E$

Observation:

Every subpath of a shortest path  
is also a shortest path

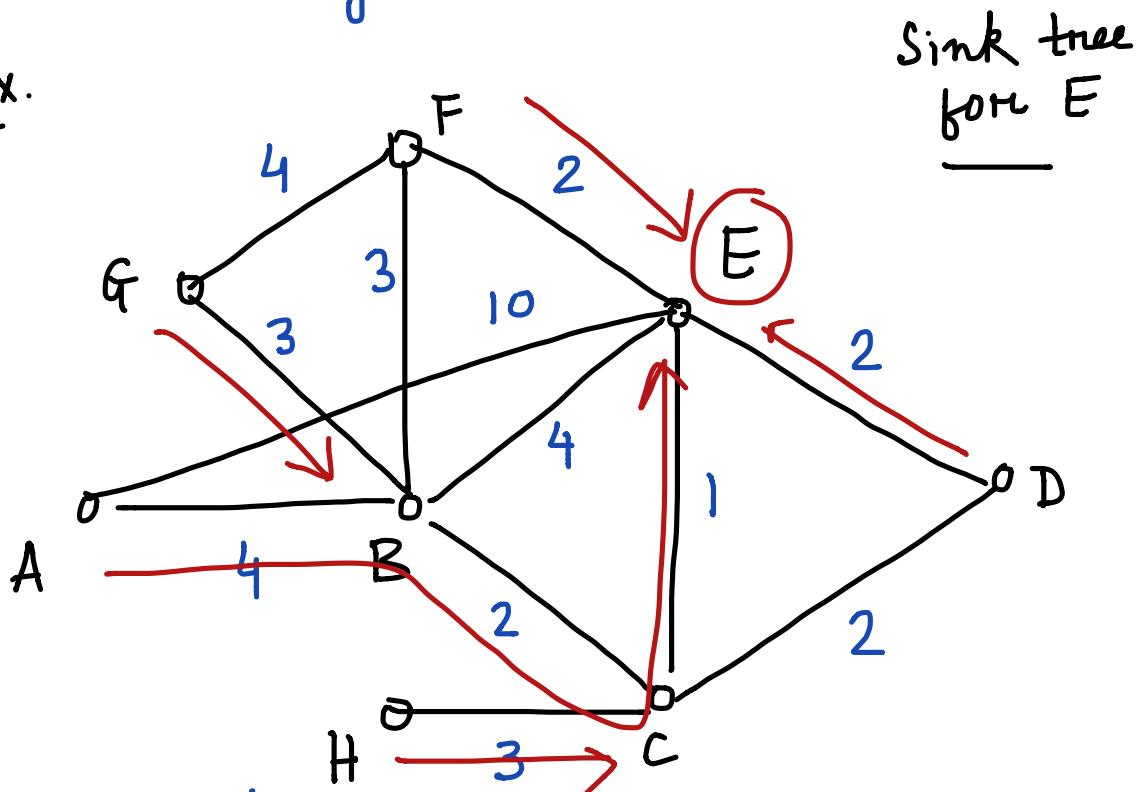
Comes as a corollary of the additive cost  
property

— Proof idea: if the subpath wasn't the  
shortest the original path wouldn't be  
shortest either.

## Sink Trees / Source Trees

- Sink tree for a destination is the union of all shortest paths towards that destination
- Similarly source tree ✓

Ex.



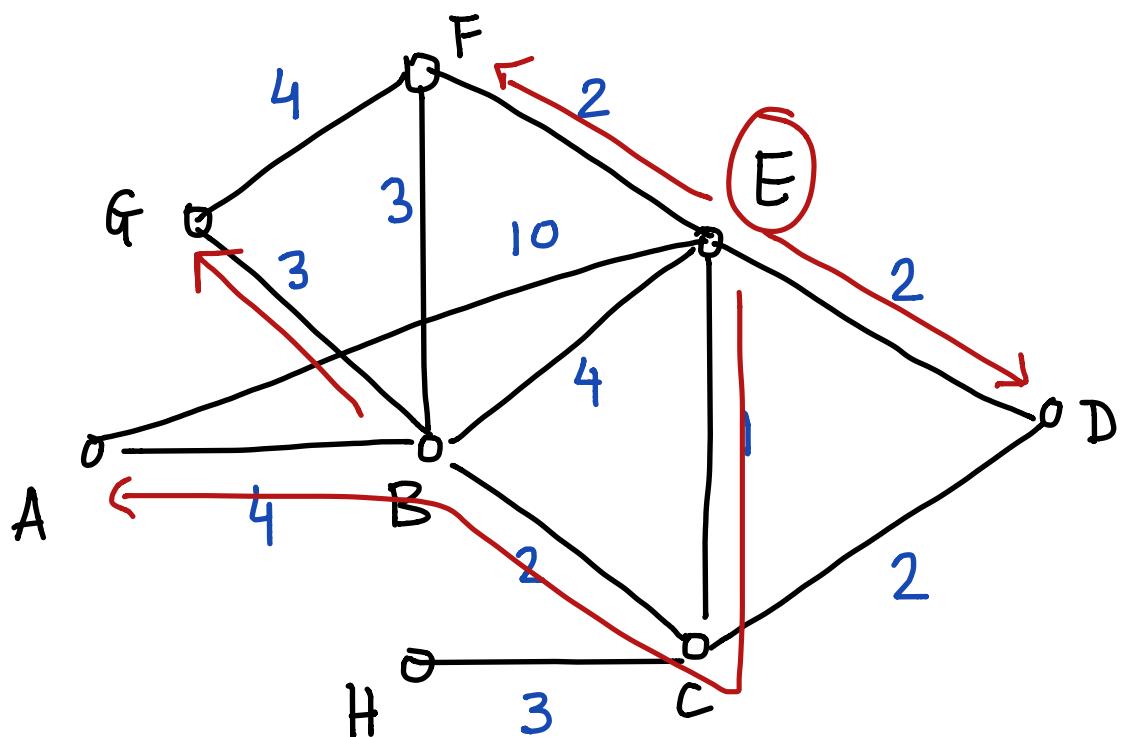
### Implications

- Only need the destination info
- Each node only needs to send to the next hop node in the sink tree
- Forwarding tables are enough.

## Algorithms to compute shortest paths

Q: How to compute shortest paths from a source to all destinations?

- Dijkstra's algorithm

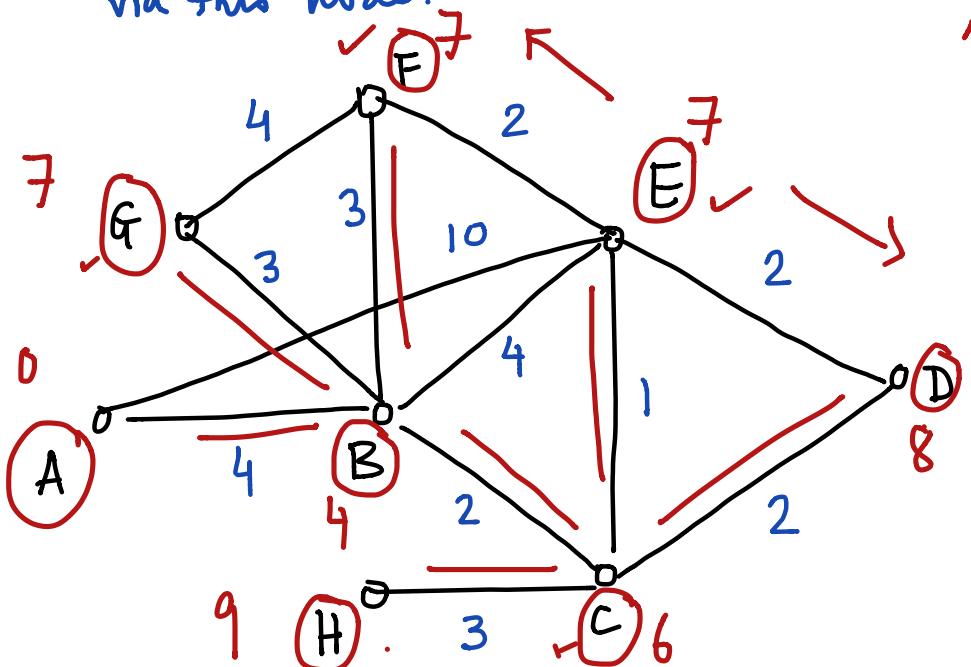


Ex: Create the forwarding table for E.

## Algorithm (Dijkstra)

① Initialization: all nodes tentative  
distance = 0 for source,  $\infty$  for all others

- ② While  $\exists$  tentative node
- find the node with lowest distance from the frozen nodes (break ties arbitrarily)
  - add link to that node into the shortest path tree
  - relax The distances of the neighbors of this newly added node, i.e., find the shortest paths to those nodes via this node.



A ✓	0	A
B	4	B
C	6	B
D	8	B
E	7	B
F	7	B
G	7	B
H	9	B

## Dijkstra endnotes

- to find the shortest paths from every node to every other node, we need to run this for different sources
- can be done carefully to minimize repetitions
- Runtime depends on the implementation of the FIND\_MIN-COST-NODE function
  - e.g., using fibonacci heap min-priority queue it is  $O(|E| + |V| \log |V|)$
- It is a centralized algorithm
- ① Does not work with negative edge weights.

Dijkstra is a greedy algorithm, fails when the costly path has an eventual negative edge later.

## Distance Vector Algorithm

✓ - simple, decentralized, distributed version of the Bellman-Ford alg.

[Bellman-Ford fixes the negative edge problem of Dijkstra, but computationally expensive]

- works, but has slow convergence among other limitations
- Link-state algorithms are used instead in practice

---

Unlike the Dijkstra setting

- ① Nodes only know the costs to their neighbors, not the topology
- ② can talk only to neighbors
- ③ can run algorithm concurrently
- ④ can fail and messages may be lost.

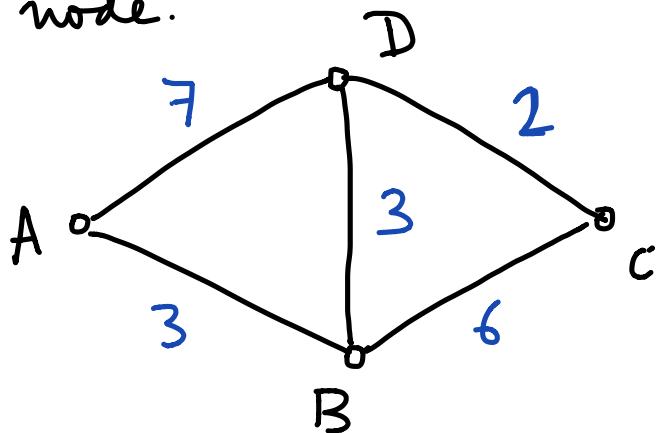
local info

## Algorithm

- Each node maintains a distance vector (distances, via nodes) to all destinations

- ① Initialize with cost = 0 to self,
  - a) to all others
- ② Send the updated distance vector to neighbors
- ③ Update the vector for each destination by selecting the min (distance from neighbors + cost to neighbors)
  - that min-cost neighbor is The via node.

E.x.

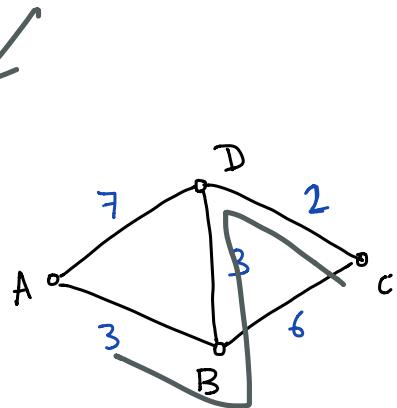


## Distance vectors

A		B		C		D	
cost	via	cost	via	cost	via	cost	via
A	0	3	A	8	D	6	B
B	3	0	B	6	B	3	B
C	8	5	D	0	C	2	C
D	7	3	D	2	D	0	D
*							

## Announcements

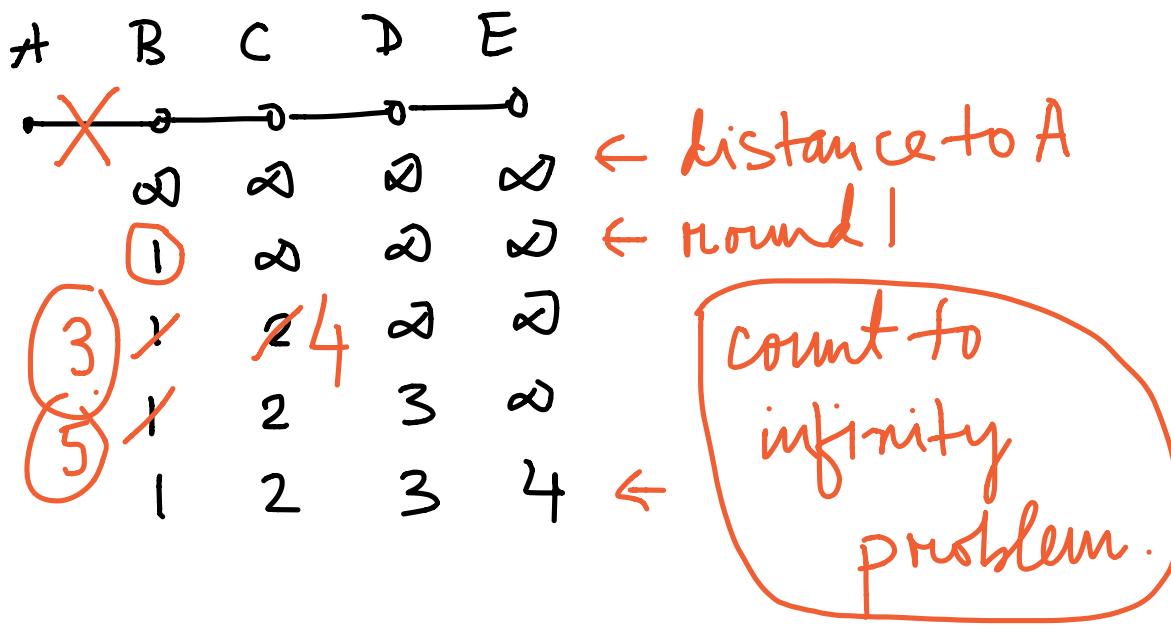
to	A	B	C	D
A	0	3	8	6
B	3	0	6	3
C	8	5	0	2
D	7	3	2	0



## Dynamics of DV

- Distance vector does local update after link failure, forgets the failed node [sets distance  $\infty$ ]
- One hop per exchange

Partitions are a problem



Heuristics are used to mitigate them  
e.g., "split horizon, poison reverse"  
don't send updates to the node where  
it is learned from

not very effective ✓

- instead link state algorithms are favored.
- 

Used in : RIP (Routing Information Protocol)