

Fair Queuing

Provides a fair quality of service (QoS)

Sharing bandwidth between flows

- Weighted Fair Queuing

FIFO:

- queue packets are stored and served in a FCFS basis.
- discard packets when queue is full.

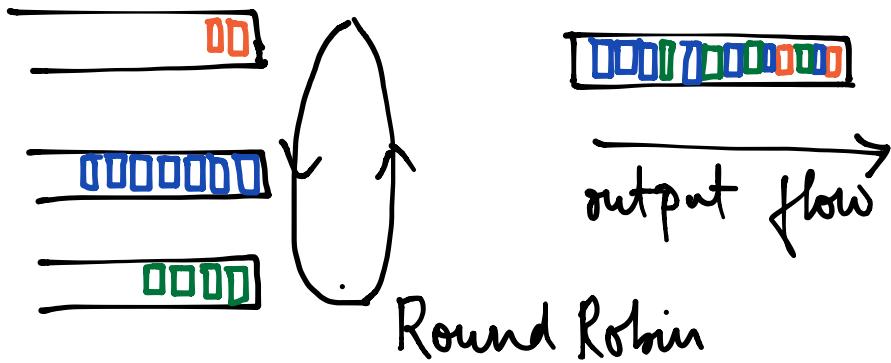


input flows

On average, a proportional representation, but may have a high variance

- difficult to maintain QoS
- malicious senders can block other flows

Round Robin Queuing



Round Robin

long-term average is still proportional,
but each flow can guarantee QoS.

The output order is no longer the
input order

- No single flow can throw out others

Problem: different packets might be of
different size. Fairness should
consider that too.

Solution: provide a bit level fairness

- Weight each of the queues s.t.
the net bit flow is equal.

1 unit per packet	\rightarrow	6 packets	}
2 units per packet	\rightarrow	3 packets	
6 units per packet	\rightarrow	1 packet	

equal
bits.

Implementations vary - but the idea is same

Weighted fair queuing is a generalization
of the above approach.

Instead of the bit rate of each flow
being equal, if we want to make it
proportional to some pre-defined ratio

E.g. if $\square \square \square = 2 : 1 : 3$ in the
previous example

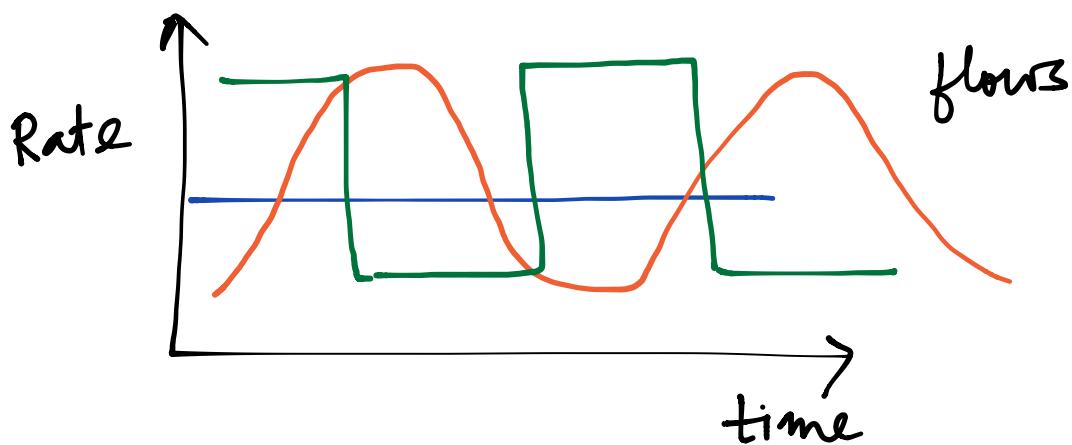
$$\begin{aligned}
 \text{Sent} &= 12 : 3 : 3 \\
 &= 12 \text{ units} : 6 \text{ units} : 18 \text{ units} \\
 &= 2 : 1 : 3
 \end{aligned}$$

Downside: difficult to implement as is

Traffic Shaping

Real traffic is bursty , so The ratio of the different flows can't be determined beforehand. The averaging has to be done on the fly.

- depends on the app , device , and stochastic network delays



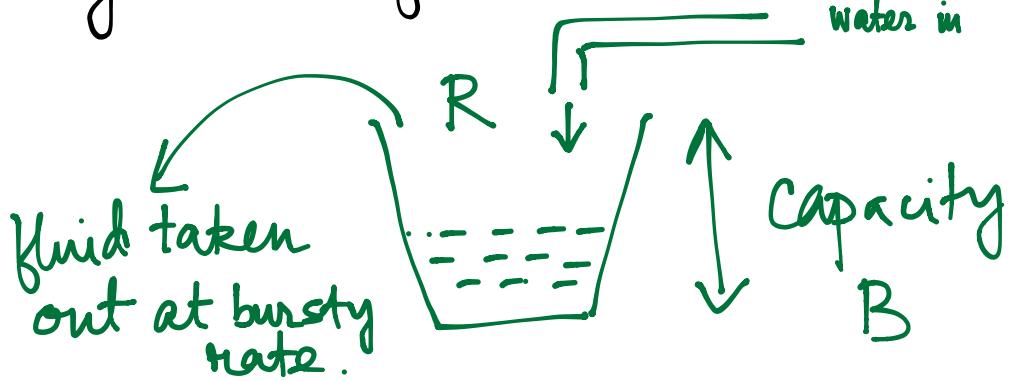
- the pattern is not regular
- not known beforehand .

TOKEN BUCKET

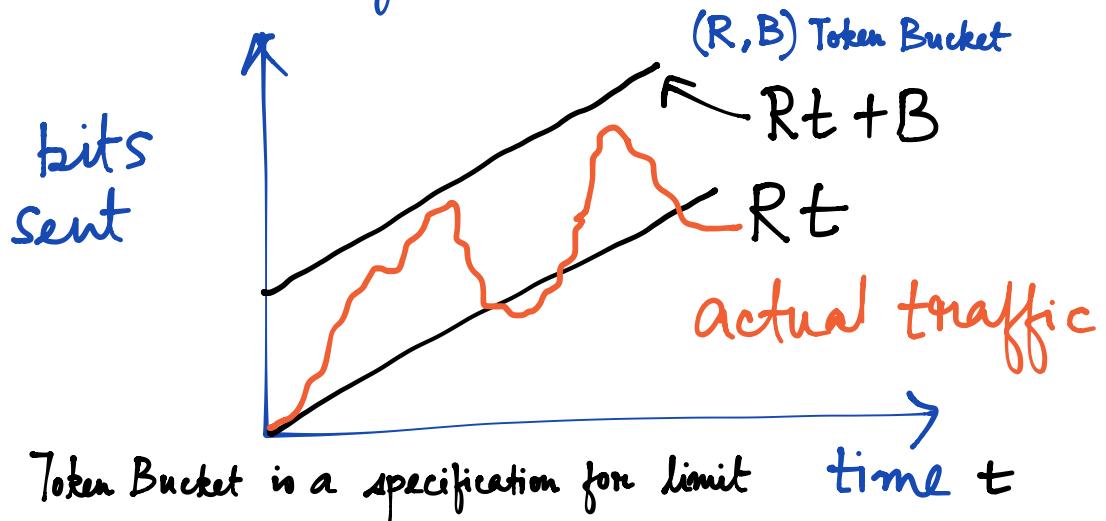
traffic shaper with two parameters

$$(R, B)$$

average rate of R bits/sec over R bursts of B bits



Time diagram



Shaping

Modifying the traffic near the source to fit within (R, B) specifications

- Run Token Bucket (R, B) at source
- Pass the tokens to the network when there are tokens automatically ensures that never more than $Rt + B$

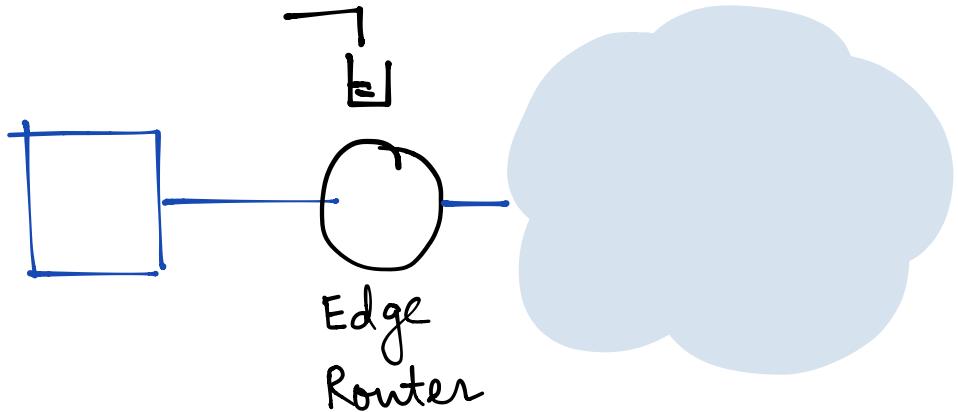
Running TB is a contract between the user and the provider.

Policing

Policing (verifies) if the traffic profile meets the (R, B) tb standard

Happens at the network

- Run (R, B) token bucket at the network edge
- Let packets into the network as long as there are tokens
- discard packets otherwise



Token bucket endnotes

- TB helps regulate the traffic
- Users have the flexibility of choosing their TB (R, B).

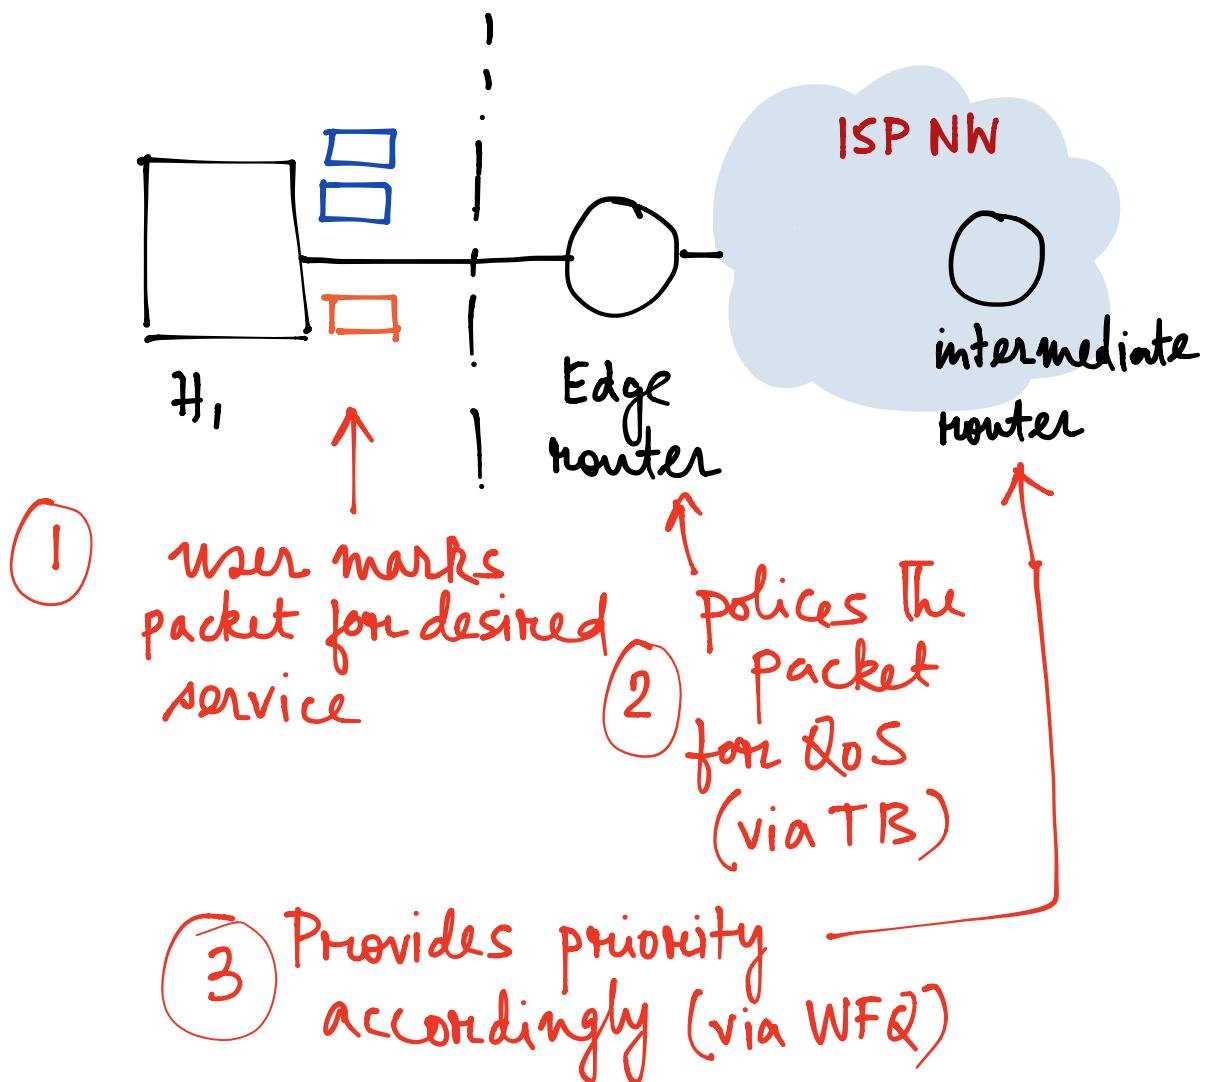
TB limits the flow, preferential treatment is given via methods like WFD

Differentiated Services

Treating different flows differently in the network.

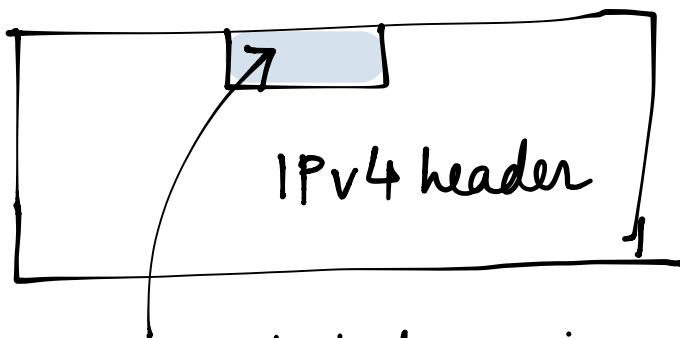
priorities can be provided by the ISP on a payment basis

Architecture of differentiated service



① Packet marking

Use bits in IPv4/IPv6 header



differentiated service
6-bit DSCP (differentiated services
code point)

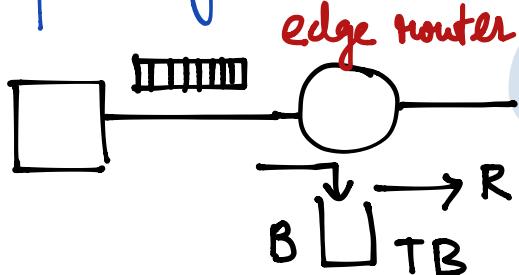
Service / meaning	DSCP	Traffic need
default forwarding/ best effort	0	elastic (P2P)
assured forwarding/ enhanced effort	10-38	average rate (streaming video)
expedited forwarding/ real time	46	low loss/delay (VOIP, gaming)
precedence/ network control	48	high priority (Routing Protocol)

- Marking packets is done by user
 - depends on network policies, e.g.,
VOIP/gaming - expedited?
 - The App can set it - can be overwritten
by the OS or router/middleboxes
 - use heuristics, e.g., port #
-

② Policing packets

ISP checks if incoming traffic meets service contract

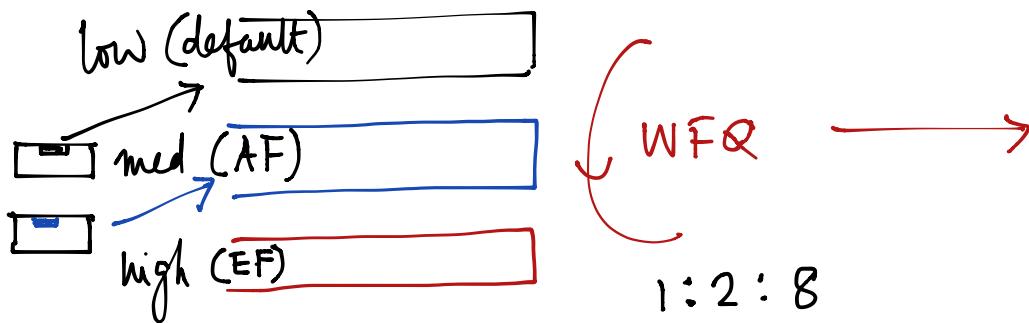
- not more expedited traffic than agreed
- only allowed markings (e.g., no network control traffic - supposed to be internal, router in ISP designs the routing protocol)
- policing is done using token bucket



Denote "out of profile" traffic
by re-marking to 0 (best
effort)

③ Forwarding Packets

- ISP use WFQ instead of FIFO
- different incoming flows (queues) correspond to the **Kinds of Service**
- DSCP values are used to map packets to the right queue



- the queue sizes for the priority classes should be reasonable - otherwise it can block all other traffic

Challenges

- QoS is an **end-to-end guarantee**
 - no benefit until both parties upgrade
- tied to **pricing** - to take care of priority overloading

Rate and Delay guarantees

So far, QoS is local (to the ISP)

- Want a guarantee across the network
 - "hard QoS" guarantee
 - e.g., like a circuit switched network
[telephone call]

Aside: TCP's connection establishment doesn't quite do this - it is between the end parties but the transmission path is not reserved
- packet losses still happen

Can we do this at all?

difficulties:

- can setup a circuit - expensive
- can we do statistical multiplexing and implement hard guarantees?

We will discuss some approaches to do this

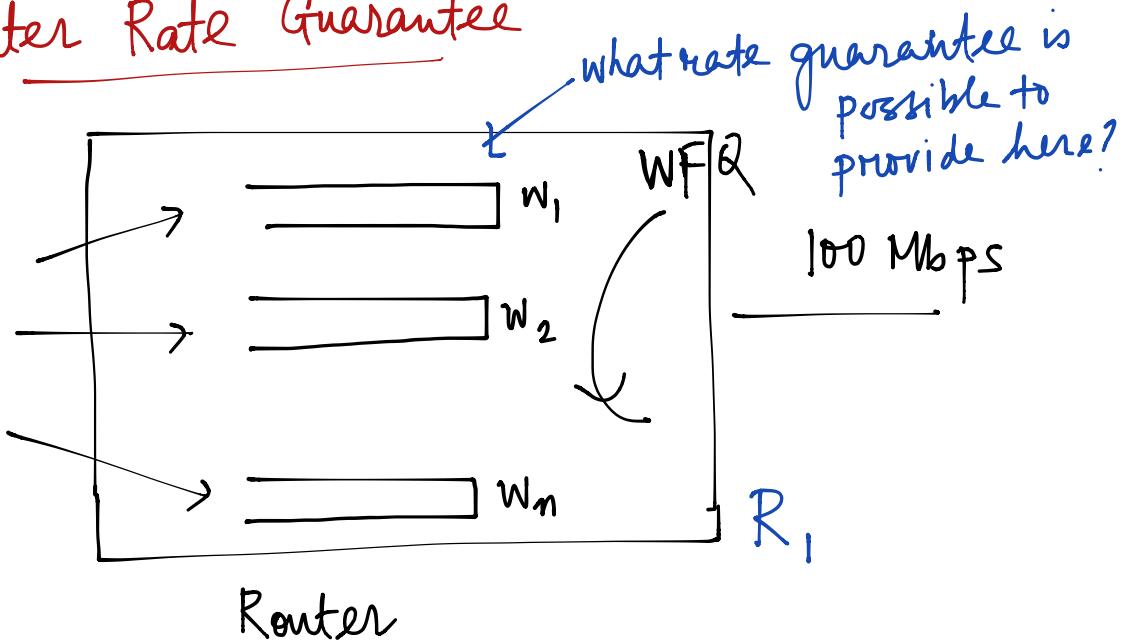
Admission Control

Requirement: a flow needs rate at least R Mbps
and delay $\leq D$ msec

decision: to admit or not? [phone call]
rejecting should be infrequent

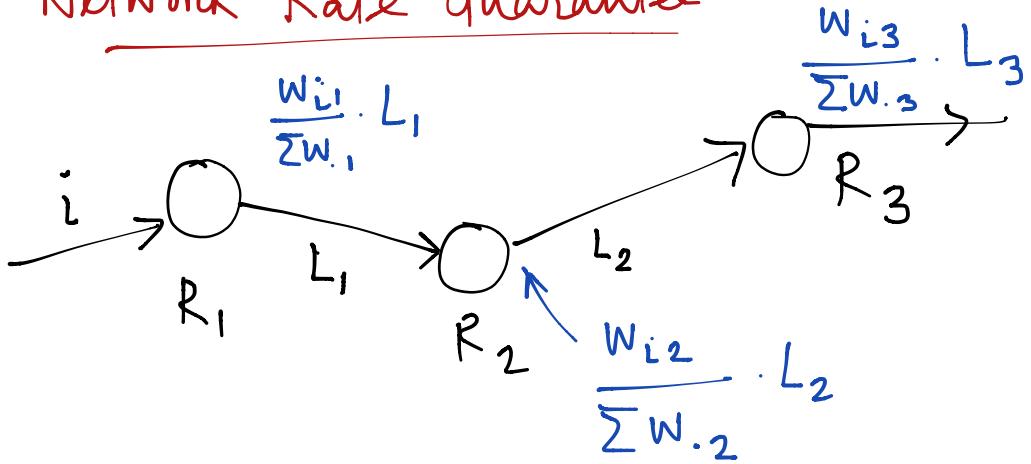
need control over admission of users to the network such that these guarantees are possible to provide.

Router Rate Guarantee



Easy to see that flow i gets $\frac{w_i}{\sum_{k=1}^n w_k} \times 100$ Mbps.

Network Rate Guarantee



The net rate is the **min** of all these rates

- to guarantee R Mbps we need

$$\min_{x \in \text{Routing path}} \frac{W_{ix} \cdot L_x}{\sum W_i \cdot x} > R$$

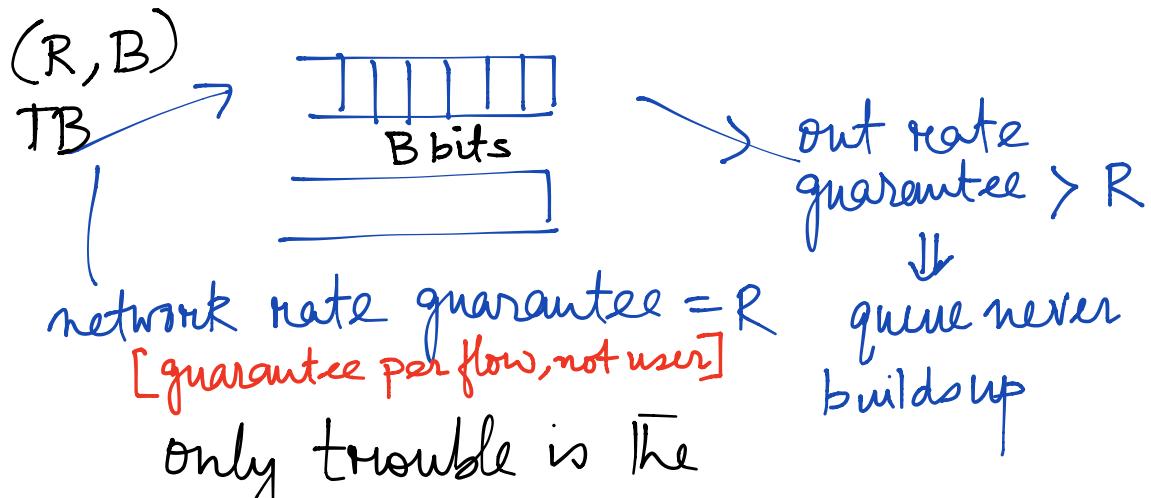
Equality has technical challenge

Delay Guarantee

- Queuing delay – transmission and propagation delays are unsurmountable
- depends on the traffic flow – if other flows get more than their minimum rate, queue

grows and delay grows

- need a max rate of all flows
 - traffic shaping (token bucket)



only trouble is the

burst — due to TB, burst can be
at most B

- this is consumed by the network min rate

$$\min_{x \in \text{Routing path}} \frac{w_i x}{\sum w \cdot x} \cdot L_x = S > R$$

- hence max delay is $\leq \frac{B}{S} < \frac{B}{R}$

Putting things together

Given : A network with

- ① (R, B) TB shaped traffic flow
- ② WFQ routers with proper weights
- ③ Sharing of output link via statistical multiplexing

Guarantee :

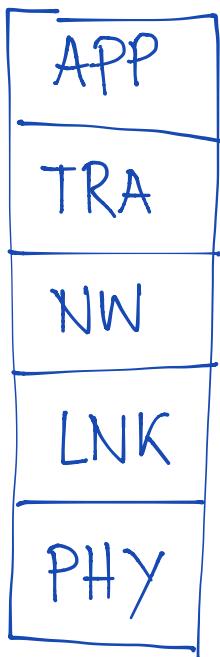
- ① minm rate = R Mbps
- ② max delay = $(t_x + \text{prop delay}) + \frac{B}{R}$

irrespective of how other flows behave
(worst case guarantee)

Summary of our learning

— A lot! can't put in a minute

— highlights



- lot detailed mechanisms drive the smooth transfer of traffic (email, WhatsApp, P2P, or Web) on the Internet
- layered structure is to help divide (distribute) tasks but the solution has to involve all

Acknowledgements

— the course TAs

— Notability App on iPad

— Xournal

— Kazam screen cast }

} linux

Video lecture
special

Stay safe &
hydrated